

Assignment-3

PythonProgramming

StudentName	Kishore G
MaximumMarks	2 Marks

Question-1:

Download the dataset: Dataset

Solution:

<https://drive.google.com/file/d/1slv-7x7CE0zAPAt0Uv-6pbO2ST2LVp5u/view>

Question-2:

Load the dataset.

Solution:

```
data=pd.read_csv("abalone.csv")
```

```
data.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

```
data.shape
```

```
(4177, 9)
```

```
data.size
```

```
37593
```

```
data.info()
```


```
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   Sex              4177 non-null  object  
1   Length           4177 non-null  float64 
2   Diameter         4177 non-null  float64 
3   Height           4177 non-null  float64 
4   Whole weight     4177 non-null  float64 
5   Shucked weight   4177 non-null  float64 
6   Viscera weight   4177 non-null  float64 
7   Shell weight     4177 non-null  float64 
8   Rings            4177 non-null  int64   
dtypes: float64(7), int64(1), object(1)
memory usage: 293.8+ KB
```

Question-3:


Perform Below Visualizations.

• Univariate Analysis


```
data.hist(figsize=(20,10), grid=False, layout=(2,4), bins=30)
plt.show()
```



```
data["Diameter"].plot(kind='hist')
```



```
data.boxplot(figsize=(20,10))  
plt.show()
```

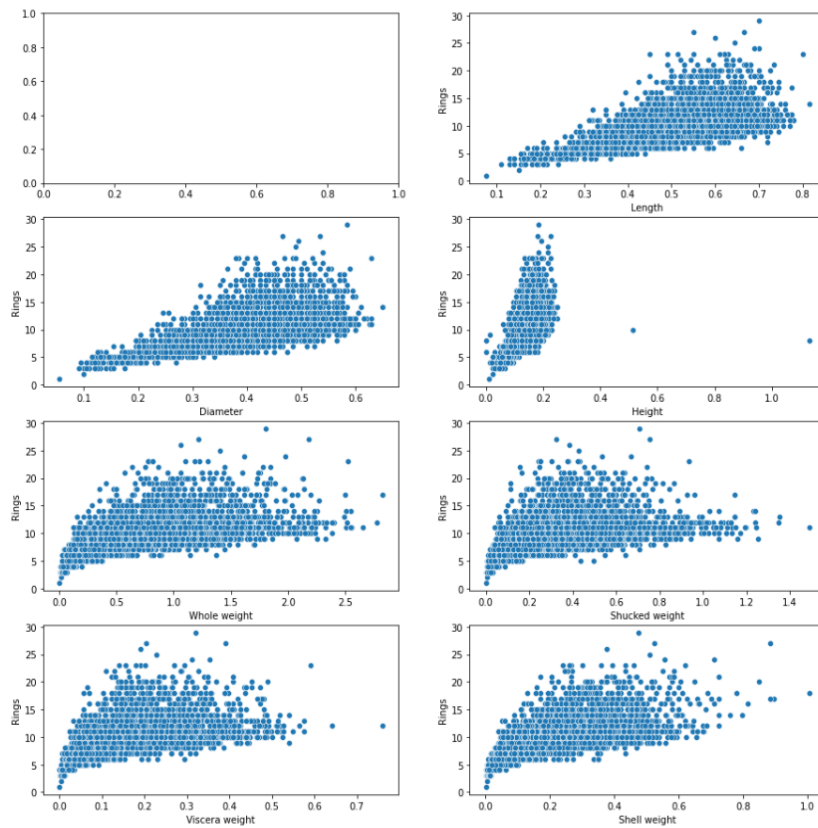


● Bi - Variate Analysis


```
data.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

```
fig, axes = plt.subplots(4,2, figsize=(15,15))
axes = axes.flatten()
for i in range(1,len(data.columns)-1):
    sns.scatterplot(x=data.iloc[:,i], y=data['Rings'], ax=axes[i])
plt.show()
```



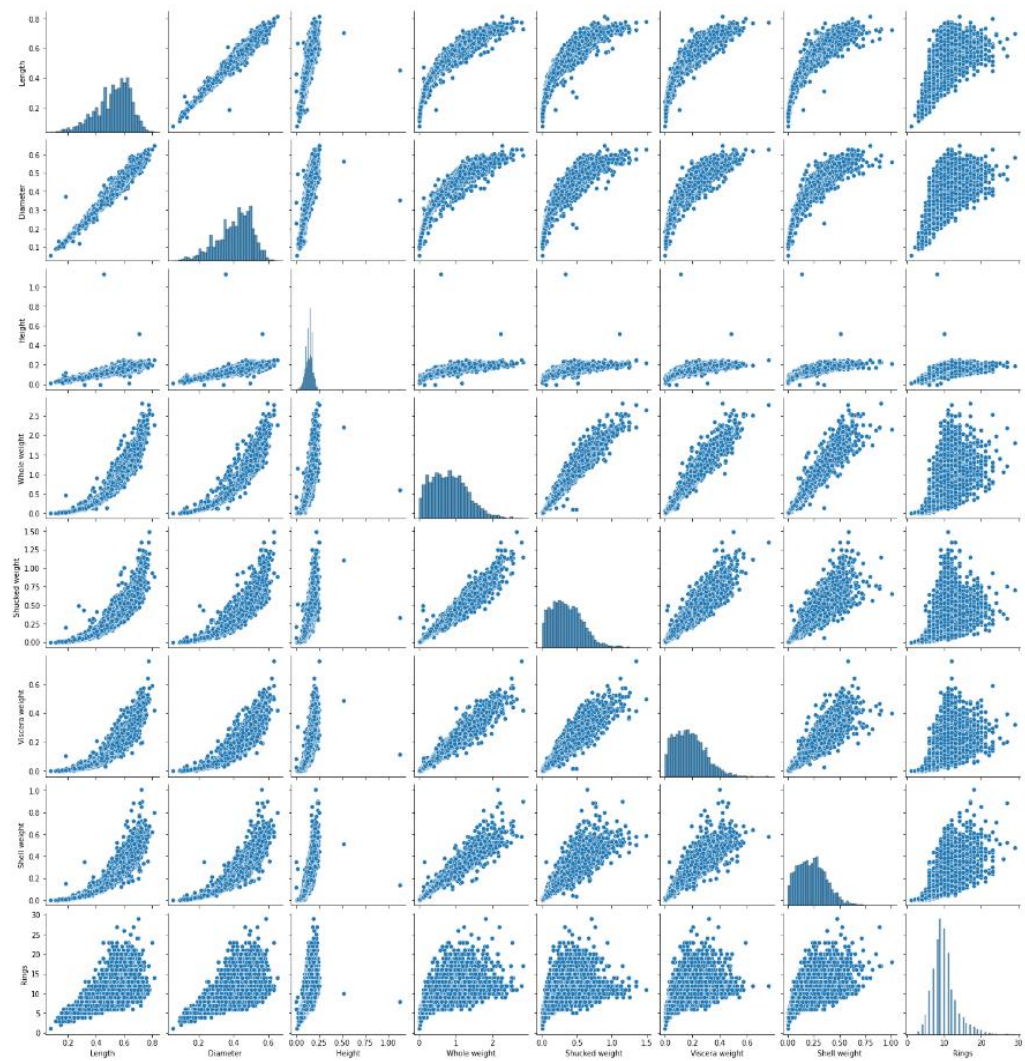
```
plt.figure(figsize=(10,5))
sns.boxenplot(y=data['Rings'], x=data['Sex'])
plt.grid()
plt.show()
data.groupby('Sex')['Rings'].describe()
```



	count	mean	std	min	25%	50%	75%	max
Sex								
F	1307.0	11.129304	3.104256	5.0	9.0	10.0	12.0	29.0
I	1342.0	7.890462	2.511554	1.0	6.0	8.0	9.0	21.0
M	1528.0	10.705497	3.026349	3.0	9.0	10.0	12.0	27.0

- Multi-VariateAnalysis

```
sns.pairplot(data)
```



Question-4:

Perform descriptive statistics on the

dataset.[Solution:](#)

```
data.describe()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	9.933684
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	1.000000
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	8.000000
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	9.000000
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	11.000000
max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	29.000000

Question-5:

Handle the Missing values.**Solution:**

```
data.isnull().any()
```

```
Sex           False
Length        False
Diameter       False
Height         False
Whole weight   False
Shucked weight False
Viscera weight False
Shell weight   False
Rings          False
dtype: bool
```

```
data.isnull().sum()
```

```
Sex           0
Length        0
Diameter       0
Height         0
Whole weight   0
Shucked weight 0
Viscera weight 0
Shell weight   0
Rings          0
dtype: int64
```


Question-6:

Find the outliers and replace the outliers**Solution:**


```
data.skew()
```

```
Length        -0.639873
Diameter       -0.609198
Height         3.128817
Whole weight   0.538959
Shucked weight 0.719098
Viscera weight 0.591852
Shell weight   0.620927
Rings          1.114102
dtype: float64
```

```
sns.boxplot(x=data['Rings'],data=data)
```



```
data['Rings'].hist()
```




```
print('skewness value of Age: ',data['Rings'].skew())
```

```
skewness value of Age: 1.114101898355677
```

```
# Flooring And Capping
```

```
Q1 = data['Rings'].quantile(0.25)
Q3 = data['Rings'].quantile(0.75)
IQR = Q3 - Q1
whisker_width = 1.5
lower_whisker = Q1 -(whisker_width*IQR)
upper_whisker = Q3 +(whisker_width*IQR)
data['Rings']=np.where(data['Rings']>upper_whisker,upper_whisker,np.where(data['Rings']<lower_whisker,lower_whisker,data['Rings']))
```

```
sns.boxplot(x=data['Rings'],data=data)
```



Question-7:

Check for Categorical columns and perform encoding.[Solution:](#)

```
data.info()
```

```
RangeIndex: 4177 entries, 0 to 4176
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Sex	4177 non-null	object
1	Length	4177 non-null	float64
2	Diameter	4177 non-null	float64
3	Height	4177 non-null	float64
4	Whole weight	4177 non-null	float64
5	Shucked weight	4177 non-null	float64
6	Viscera weight	4177 non-null	float64
7	Shell weight	4177 non-null	float64
8	Rings	4177 non-null	float64

```
dtypes: float64(8), object(1)
```

```
memory usage: 293.8+ KB
```

```
#Label Encoding
```

```
from sklearn.preprocessing import LabelEncoder
```

```
le=LabelEncoder()
```

```
data['Sex']=le.fit_transform(data['Sex'])
```

```
data.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15.0
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7.0
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9.0
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10.0
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7.0

```
data["Sex"].unique()
```

```
array([2, 0, 1])
```

Question-8:

Split the data into dependent and independent variables.**Solution:**

```
data.head(5)
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15.0
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7.0
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9.0
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10.0
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7.0

```
X = data.iloc[:, 0:7]
Y = data.iloc[:, -1]
```

X

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395
...
4172	0	0.565	0.450	0.165	0.8870	0.3700	0.2390
4173	2	0.590	0.440	0.135	0.9660	0.4390	0.2145
4174	2	0.600	0.475	0.205	1.1760	0.5255	0.2875
4175	0	0.625	0.485	0.150	1.0945	0.5310	0.2610
4176	2	0.710	0.555	0.195	1.9485	0.9455	0.3765

4177 rows × 7 columns

Y

```
0    15.0
1     7.0
2     9.0
3    10.0
4     7.0
...
4172  11.0
4173  10.0
4174   9.0
4175  10.0
4176  12.0
Name: Rings, Length: 4177, dtype: float64
```

```
y.shape
```

```
(4177, 0)
```

Question-9:

Scale the independent variables.**Solution:**

```
from sklearn.preprocessing import StandardScaler
ss = StandardScaler()
X_scaled = ss.fit_transform(X)
```


Question-10:

Split the data into training and testing

Solution:

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X_scaled, Y, test_size = 0.3, random_state = 1)
```

Question-11:

Build the Model

Training the Model Testing the Model

```
import csv
with open("abalone.csv") as csv_file:
    csv_reader = csv.reader(csv_file)
    data = pd.DataFrame([csv_reader], index = None)
for val in list(data[1]):
    print(val)
```

```
['M', '0.455', '0.365', '0.095', '0.514', '0.2245', '0.101', '0.15', '15']
```

Question-12:

Train the Model and Test the Model

1.Linear Regression 2.Ridge 3.Decision Tree Regression 4.KNeighborsRegressor

```
#importing all the necessary models and metrics
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor

from sklearn.metrics import mean_squared_error, r2_score
```

1. Linear Regression

```
lr = LinearRegression()
lr.fit(x_train, y_train)
```

LinearRegression()

```
#Testing the model
lr_test_pred = lr.predict(x_test)
```

lr_test_pred

```
array([8.49722433, 7.64369059, 7.82520883, ..., 8.55677832, 9.02884473,
       5.96561877])
```

```
#measuring the performance
mse = mean_squared_error(y_test, lr_test_pred)
print('Mean Squared error of testing Set: %2f'%mse)
```

Mean Squared error of testing Set: 3.524602

```
p = r2_score(y_test, lr_test_pred)
print('R2 Score of testing set:%.2f'%p)
```

R2 Score of testing set:0.52

2. Ridge

```
ridge_mod = Ridge(alpha=0.01, normalize=True)
ridge_mod.fit(x_train, y_train)
ridge_mod.fit(x_test, y_test)
```

Ridge(alpha=0.01, normalize=True)

```
#Testing the model
ridge_model_pred = ridge_mod.predict(x_test)
```

ridge_model_pred

```
array([8.54031033, 8.48463396, 7.96838487, ..., 8.77493484, 9.03881023,
       5.83582085])
```

```
#Measuring the performance
acc = r2_score(y_test, ridge_model_pred)
print('Score of testing Set: %2f'%acc)
```

Score of testing Set: 0.523227

3. Decision Tree Regression

```
dt = DecisionTreeRegressor()
dt.fit(x_train, y_train)
```

DecisionTreeRegressor()

```
#Testing the model
dt_test_pred = dt.predict(x_test)
```

dt_test_pred

array([12., 9., 10., ..., 7., 9., 4.])

Question-13:

Measure the performance using Metrics.

```
#Measuring the Performance
dacc = mean_squared_error(y_test, dt_test_pred)
print('Mean Squared Error of testing Set: %2f'%dacc)
```

Mean Squared Error of testing Set: 6.126994

4. KNN Regression

```
knn = KNeighborsRegressor(n_neighbors = 4 )
knn.fit(x_train, y_train)
knn.fit(x_test, y_test)
```

KNeighborsRegressor(n_neighbors=4)

```
#Testing the Model
knn_test_pred = knn.predict(x_test)
```

knn_test_pred

array([8.75, 9.5 , 10.5 , ..., 8. , 7.5 , 5.])

```
#Measuring the Performance
kacc= r2_score(knn_test_pred,y_test)
print('Score of testing Set: %2f'%kacc)
```

Score of testing Set: 0.400555

```
kmse = mean_squared_error(knn_test_pred,y_test)
print('Score of testing Set: %2f'%kmse)
```

Score of testing Set: 2.602460